

DELFT UNIVERSITY OF TECHNOLOGY

BACHELOR GRADUATION PROJECT

INITIAL RESEARCH REPORT

UrbanSearch

Authors:

Tom BRUNNER
Marko MALIŠ
Gijs REICHERT
Piet VAN AGTMAAL

Supervisor:

Claudia HAUFF

Clients:

Evert MEIJERS
Antoine PERIS

April 28, 2017

Abstract

It is yet to be discovered how the importance of cities in the global network can be elucidated. In this document, the requirements for a solution to this problem are identified, as well as possible issues that might arise. We explore various methods and select the most appropriate ones that, if combined, satisfy these requirements and tackle the identified issues, in preparation for an attempt at solving the problem.

Keywords: urban, network, data mining, document analysis, filtering

Contents

1	Introduction	3
2	Related Work	3
3	Requirement Analysis	3
3.1	Design Goals	3
3.2	Product Requirements	4
3.2.1	Must Have	4
3.2.2	Should Have	4
3.2.3	Could Have	4
3.2.4	Would Like	5
4	Framework and Tools	5
4.1	Gathering the Data	5
4.1.1	Common Crawl	5
4.1.2	Delpher	6
4.1.3	Eurostat	6
4.2	Data Storage and Ingestion	6
4.2.1	Elasticsearch	6
4.2.2	Hadoop	7
4.2.3	MrJob	7
4.2.4	Neo4j	7
4.3	Filtering and Classifying	7
4.3.1	Clustering	7
4.3.2	Filtering	7
4.3.3	Machine Learning for classification	7
4.3.4	TF-IDF	8
4.4	Search Queries	9
4.4.1	Enter Queries	9
4.4.2	Get Results	9
4.4.3	Specifications	9
4.5	Visualisation	9
4.5.1	neo4j?	9
4.5.2	Connection between cities	9
4.5.3	The Strength of these connections	9
5	Conclusion	9

1 Introduction

Common belief is that agglomeration benefits are key to economic growth [15]. However, it remains uncertain whether it is these benefits that primarily cause economic growth, or an increase in the position of the agglomeration in the global urban network.

The huge amount of textual data generated online and the numerous historic archives are great sources of information on social and economic behaviours. They constitute the bulk of information flowing among each other. Advanced text mining on newspapers and web pages containing city names would allow for a better understanding of the role of information in shaping urban systems. Similar to research efforts in other domains, such as financial trade [16] and sales forecasting [20], the idea is to develop search queries that capture urban-urban interactions. These interactions are retrieved from information corpora through the co-occurrence of geographical names in textual data. An example of such a query on the Google search engine¹ is "Rotterdam + Amsterdam " OR "Amsterdam + Rotterdam", which searches for the co-occurrence of Amsterdam and Rotterdam. We thus answer the following question **FIXME: correct?**: what approach is best suited for extracting and visualising the global urban network from open data? **FIXME: define subquestions?**

First, we discuss related work in section 2. Second, we identify the requirements for a solution to the problem and discuss issues that might arise in section 3. Third, we develop a methodology for a framework that satisfies the requirements and tackles the issues in section 4. Last, we conclude in section 5 with the results of our research.

2 Related Work

3 Requirement Analysis

In this section, we first define the design goals. Then, we list the requirements which the **FIXME thingy** meet. We use the MoSCoW method [2] as a prioritisation technique. Four levels of priority are defined: must have, should have, could have and won't have (also known as would like).

3.1 Design Goals

The high-level design goals for this project have been provided by the client. These serve as a guideline to determine the priority label of the specific requirements, as defined in section 3.2 and are listed here, ordered by priority.

credible The client wants to impugn a widely spread belief. Therefore, the basis on which he does that must be sound.

understandable The results of the **FIXME system?** should be visually understandable, so it is easy for the client to deduce conclusions from them. Additionally, retrievable numeric data enable the client to further investigate the results outside of the scope of the **FIXME system?**, should the need arise.

scalable The client has expressed his concerns that restricting the **FIXME system?** to a set of non-English domains might impair the probability that his research will be published in an acknowledged journal. Allowing for investigating other domains would greatly help the client in a later stadium.

¹<https://www.google.com>

pluggable The client conveyed it might be interesting to let the **FIXME system?** perform analysis on different data sets without the need of a developer.

3.2 Product Requirements

3.2.1 Must Have

Requirements labelled as "must have" are key to the minimal performance of the **FIXME product/whatever**. If they are not met, the **FIXME: thingy** can be considered a failure.

1. General
 - (a) Adding city names
 - (b) Grouping relations and "zooming" on these relations
2. Search Engine
 - (a) Filter results
 - (b) Data mining
3. Filtering
 - (a) Logic Filters
 - (b) Relations Filters
4. Machine Learning
 - (a) Types of relations
5. Visualisation
 - (a) Statistics of relations? Query relations
 - (b) Strength of relations
 - (c) Types: ML CBS defined

3.2.2 Should Have

"Should have" requirements are those that greatly improve system performance and/or usability but might not fit in the available development time.

1. General
 - (a) Pluggable datasets
2. Machine Learning
 - (a) Generalising relations, grouping relations

3.2.3 Could Have

Requirements labelled as "could have" are useful and should be included in the system if time and resources permit.

1. General
 - (a) International city names
2. Visualisation
 - (a) Front end for the app

3.2.4 Would Like

”Would like” requirements have been agreed upon as not important to deliver within the current time schedule. However, they can be included in future releases.

4 Framework and Tools

In this section, we evaluate what data sources are useful. Additionally, we discuss several methods and tools that can be helpful in storing and ingesting data. Furthermore, we describe numerous methods to filter and classify textual data. Then, we elaborate on different methods to perform queries with. We conclude with an overview of the available visualisation tools that can be used for displaying the results of the analysis.

4.1 Gathering the Data

4.1.1 Common Crawl

Common Crawl [3] is a freely accessible corpus of the pages across the web. Their data are updated and released on a monthly basis. Many researchers have used the data for varying purposes [18] [11] [17]. Since the project requires us to crawl the web (see section **FIXME**), the corpus is a very suitable candidate for us to work with.

The data of Common Crawl come in three formats²:

WARC This is the default and most verbose format. It stores the HTTP-response, information about the request and meta-data on the crawl process itself. The content is stored as HTML-content.

WAT Files of this type contain important meta-data, such as link addresses, about the WARC-records. This meta-data is computed for each of the three types of records (meta-data, request, and response). The textual content of the page is not present in this format.

WET This format only contains extracted plain text. No HTML tags are present in this text. For our purposes, this is the most useful format.

For extracting data from Common Crawl, many open-source libraries are available. Common Crawl’s official website refers to **cdx-index-client**³ as a command line interface to their data. It allows for, among others, specifying which dataset to use, supports multiple output formats (plain text, gzip or JSON) and can run in parallel.

A simple query on the latest index using the online interface⁴ yields 1676 pages of 15000 entries each, which are roughly 25 million entries in total. However, there are over 5.5 million registered domain names with top level domain **.nl**⁵. One would expect many more pages to exist with that number of domains. There are several explanations for this, including:

- Common large websites, such as **www.google.nl** and **www.wikipedia.nl** have not been fully indexed by Common Crawl, because their ”parents”, **www.google.com** and **www.wikipedia.org** have already been indexed almost entirely.
- Not every website allows their pages to be crawled. According to Common Crawl’s official website, their bots can be blocked via the common **robots.txt**. Additionally, they honour so-called **no-follow** attributes that prevents the crawler to follow embedded links. Sites that use these features are therefore partially or not at all included in the indices of Common Crawl.

²<https://gist.github.com/Smerity/e750f0ef0ab9aa366558>

³<https://github.com/ikreymer/cdx-index-client>

⁴http://index.commoncrawl.org/CC-MAIN-2017-13-index?url=*.nl&output=json&showNumPages=true

⁵https://www.sidn.nl/a/knowledge-and-development/statistics?language_id=2

4.1.2 Delpher

Delpher [1] is a library consisting of millions of Dutch digitalised newspapers, books and magazines of ages differing from the fifteenth century up until now. Their data are free to use for private purposes. All data can be accessed however, if the organisation doing research explicitly asks for permission to use the data and agrees to Delpher's terms. Delpher could be useful to perform historical analysis and see trends over longer periods of time. This cannot be done using web pages as source, as the Internet is still very young.

Should historical analysis be desired, Delpher is the data source to use. Since their data are texts from reliable sources, there is no noise to be filtered out. Additionally, some valuable information can be extracted, such as area of distribution, cross-location referencing and target audience. These aid in the process of measuring the strengths of relationships.

4.1.3 Eurostat

Eurostat is the statistical office of the European Union situated in Luxembourg. Its mission is to provide high quality statistics for Europe [6]. We identified Eurostat as a source that is not useful for this particular problem. Although Eurostat contains a lot of statistics on European cities, there is not enough useful information which contributes to giving more insight into the network connectivity of cities. Therefore, we did not include Eurostat as an information source.

4.2 Data Storage and Ingestion

4.2.1 Elasticsearch

Elasticsearch is an open-source search engine which centrally stores your data [4]. It is a fast and scalable solution that was designed with big data search in mind. According to Kononenko et al. [9] Elasticsearch has some significant advantages in comparison with traditional relational databases. Two of these advantages are scalability and performance.

Scalability According to Elasticsearch [4] their product has no problem with scaling horizontally. It automatically manages indices and queries distributed across a cluster. This is an important feature as it is likely that the amount of data that our solution will use and process will increase and we do not want to keep upgrading the server that contains database, which would be vertical scalability. **FIXME: Referentie naar dat we alleen NL doen nu? En daarom scalability nogal belangrijk is**

Performance Because Elasticsearch was designed to handle documents and perform full-text search it not surprising it performs well doing this. As we're going to be using the same kind of input data we expect Elasticsearch as the most choice. Kononenko et al. found that while scalability and schema-free documents are common for NoSQL systems, the combination of all three (scalability, agility, and performance) in one system is what makes Elasticsearch stand out from other systems. Following this, we conclude Elasticsearch would be a good choice as a data storage and search platform for our **FIXME: solution—product—project.. (We moeten even kiezen welke we aanhouden zodat we dat overal hetzelfde doen.**

Downsides A downside of Elasticsearch is that it does not have any form of security out of the box. This means that everyone with the server address could access the data. This is not a problem for CommonCrawl data, as this was already available online anyway. However, when using for example Delpher or other sources you need a license this becomes a problem. Next to that, it would also be possible for anyone to meddle with the data in Elasticsearch making the data unreliable. Elasticsearch provides a Security package for which you unfortunately need a

paid license. However, to secure Elasticsearch while making it available for users we could use a plugin such as Search Guard⁶ or use a special proxy as proposed by Kononenko et al. [9].

Another downside of Elasticsearch is that transactions involving multiple documents⁷ are not ACIDic. Where ACID stands for the four properties atomicity, consistency, isolation and durability regarding transactions in database systems [7]. This means that we need to keep concurrency problems in mind and will probably need to enable some locking to prevent these concurrency problems when performing transactions on multiple documents.

4.2.2 Hadoop

Because we are designing a

4.2.3 MrJob

4.2.4 Neo4j

Neo4j is a highly scalable native graph database that leverages data relationships as first-class entities [12], enabling enterprises of any size to connect their data and use the relationships to improve their businesses. It is the single highly-scalable, fast and ACID compliant graph database available. Additionally, it is free to use for non-commercial applications.

There are some specific aspects of Neo4j that make it a very suitable candidate for the **FIXME project**. These are: annotations, relations and properties. For our purposes, it can be used to store the global network of urban relations and visualise it through extracting the graph from the database. Since data already form a graph, this is an obvious task.

Scalability As stated on their official website, Neo4j is "the only enterprise-strength graph database that combines native graph storage, scalable architecture optimised for speed, and ACID compliance (see section 4.2.1 for a short explanation). To illustrate how scalable Neo4j is, consider that very large companies such as ebay, Cisco, Walmart, HP and LinkedIn⁸ use it in their mission-critical systems.

Performance Holzschuher and Peinl compared the performance of Neo4j to the more classic and commonly used NoSQL databases and found out that the more natural representation of relationships resulted in significant performance increase gains [8]. TODO: expand

4.3 Filtering and Classifying

4.3.1 Clustering

4.3.2 Filtering

4.3.3 Machine Learning for classification

Text based machine learning can be done in two different ways depending on whether the text is structured or unstructured. Text is structured when sentences are used, meaning grammar is used. However on websites you also often find loose words. In the case of structured texts a word may say something about the next word in the sentence, therefor other techniques can be used (for example n-grams) then in case of unstructured texts. Machine learning for website classification works by using the raw text from a website. Because the text from websites is often unstructured the 'bag of words' model is used. This model counts how often each word is used. There are 2 libraries available which contain most steps needed. There is scikit-learn [10] and TensorFlow [19] for Python and Weka [13] for Java. The machine learning works in 4 steps:

⁶<http://floragunn.com/searchguard/>

⁷<https://www.elastic.co/guide/en/elasticsearch/guide/current/concurrency-solutions.html>

⁸<https://neo4j.com/customers/>

1. Creating a feature extractor

Given text from a website, returns the "features" from this text. Features are the words that occur in the text and the number of occurrences. Before this data is extracted stop words (the, is, at etc) are removed and the rest of the words are stemmed meaning all words will be changed to their root-forms (features - feature, controlled - control). The program Snowball [14] and NLKT [5] (which uses the snowball version) have a dutch implementation for this, although it might need to be improved a bit.

Afterwards, to prepare the features for the machine learning algorithm. We need to give each feature a numeric id. Count each of these tokens. And we need to normalise the tokens. For this scikit-learn provides algorithms.

2. Manually labelling

For each of the classes (e.g. business, tourism, art etc) we select a few websites we know fit to that category. From these websites all the words will be extracted and their occurrence will be counted. Possibly some normalisation functions are applied to get better values. We call these values the weight for each word for each class. From this we create a two dimensional array with in the rows each of the classes and in the columns all different words. We fill the fields with the weights or a zero if the words don't occur.

3. Generating a classifier

The array is fed to a learning algorithms. This will generate a classifier. Several classifiers (provided by scikit-learn and TensorFlow) are:

- (a) LDA & QDA
- (b) SVM - LibSVM ($O(n^3?)$)/*LibLinear*
- (c) Stochastic Gradient Descent
- (d) K-Nearest Neighbour
- (e) GPR
- (f) Decision trees (usually C4.5)
- (g) Ensemble Methods
- (h) MultiClass and MultiLabel methods
- (i) Neural Networks

4. Entering new examples

When a new (unlabelled) example (website) comes - extract the features and feed it to your classifier - it will tell you what it thinks it is (and usually - what is the probability the classifier is correct). Afterwards the classifier is updated to include new features extracted from the example.

todos:
n-grams / shingling

4.3.4 TF-IDF

basic idea: 1. using training data to assign values on words - filter meaningless words - assign words with highest value as categories? 2. Do the same on training data for each category (choose a few documents manually per category) and then check for websites for which categories has the highest value.

4.4 Search Queries

4.4.1 Enter Queries

4.4.2 Get Results

4.4.3 Specifications

4.5 Visualisation

4.5.1 neo4j?

4.5.2 Connection between cities

4.5.3 The Strength of these connections

5 Conclusion

References

- [1] Koninklijke Bibliotheek. Delpher - boeken kranten tijdschriften. <http://delpher.nl/>. Accessed: 2017-04-26.
- [2] Dai Clegg and Richard Barker. *Case method fast-track: a RAD approach*. Addison-Wesley Longman Publishing Co., Inc., 1994.
- [3] Common Crawl. Common crawl. <https://commoncrawl.org/>, 2017. Accessed: 2017-04-25.
- [4] Elasticsearch. Elasticsearch. <https://www.elastic.co/products/elasticsearch>. Accessed: 2017-04-26.
- [5] Alex Rudnick et al. Nlkt. <http://www.nltk.org/api/nltk.stem.html>, 2017.
- [6] Eurostat: About. <http://ec.europa.eu/eurostat/about/overview>. Accessed: 2017-04-26.
- [7] Theo Haerder and Andreas Reuter. Principles of transaction-oriented database recovery. *ACM Computing Surveys (CSUR)*, 15(4):287–317, 1983.
- [8] Florian Holzschuher and René Peinl. Performance of graph query languages: comparison of cypher, gremlin and native access in neo4j. In *Proceedings of the Joint EDBT/ICDT 2013 Workshops*, pages 195–204. ACM, 2013.
- [9] Oleksii Kononenko, Olga Baysal, Reid Holmes, and Michael W. Godfrey. Mining modern repositories with elasticsearch. In *Proceedings of the 11th Working Conference on Mining Software Repositories*, MSR 2014, pages 328–331, New York, NY, USA, 2014. ACM.
- [10] Scikit learn developers. Scikit-learn. <http://scikit-learn.org/stable/index.html>, 2017. Accessed: 2017-04-26].
- [11] Hannes Mühleisen and Christian Bizer. Web data commons-extracting structured data from two large web corpora. *LDOW*, 937:133–145, 2012.
- [12] Neo4j. Neo4j, the world’s leading graph database. <https://www.neo4j.com>. Accessed: 2017-04-26.
- [13] University of Waikato. Weka. <http://www.cs.waikato.ac.nz/ml/weka/>, 2017. Accessed: 2017-04-26.
- [14] Richard Boulton Olly Betts. Dsnowball. <http://snowball.tartarus.org/algorithms/dutch/stemmer.html>, 2017.

- [15] Michael E Porter. Location, competition, and economic development: Local clusters in a global economy. *Economic development quarterly*, 14(1):15–34, 2000.
- [16] Tobias Preis, Helen Susannah Moat, and H Eugene Stanley. Quantifying trading behavior in financial markets using google trends. *Nature: scientific reports*, 2013.
- [17] Sameer Singh, Amarnag Subramanya, Fernando Pereira, and Andrew McCallum. Wikilinks: A large-scale cross-document coreference corpus labeled via links to wikipedia. *University of Massachusetts, Amherst, Tech. Rep. UM-CS-2012-015*, 2012.
- [18] Jason R Smith, Herve Saint-Amand, Magdalena Plamada, Philipp Koehn, Chris Callison-Burch, and Adam Lopez. Dirt cheap web-scale parallel text from the common crawl. In *ACL (1)*, pages 1374–1383, 2013.
- [19] Google Brain Team. Tensorflow. <https://www.tensorflow.org/>, 2017. Accessed: 2017-04-28].
- [20] Lynn Wu and Erik Brynjolfsson. The future of prediction: How google searches foreshadow housing prices and sales. In *Economic analysis of the digital economy*, pages 89–118. University of Chicago Press, 2014.