

DELFT UNIVERSITY OF TECHNOLOGY

BACHELOR GRADUATION PROJECT

INITIAL RESEARCH REPORT

UrbanSearch

Authors:

Tom BRUNNER

Marko MALIŠ

Gijs REICHERT

Piet VAN AGTMAAL

Supervisor:

Claudia HAUFF

Clients:

Evert MEIJERS

Antoine PERIS

June 19, 2017

Abstract

It is hard if not impossible to measure the strength of relationships between cities using existing technologies. Because of this, it remains uncertain how exactly economic growth is affected by urbanisation. The universally accepted explanation is that only increasing size of cities affects economic growth [20]. In this report, we develop a methodology that allows for determining intercity relationship strengths, using open data. For this, we evaluate whether graph databases like Neo4j [17] or document search engines such as Elasticsearch [8] are best suited and describe machine learning algorithms for categorising data based on the co-occurrence of city names. Additionally, we present visualisation techniques to be able to intuitively analyse the results.

Keywords: urban, relation strength, data mining, graph, classification, filtering

Contents

| | | |
|----------|---|-----------|
| 1 | Introduction | 3 |
| 2 | Related Work | 3 |
| 3 | Requirement Analysis | 4 |
| 3.1 | User Stories | 4 |
| 3.2 | Design Goals | 4 |
| 3.3 | Product Requirements | 5 |
| 3.3.1 | Must Have | 5 |
| 3.3.2 | Should Have | 5 |
| 3.3.3 | Could Have | 6 |
| 3.3.4 | Would Like | 6 |
| 3.4 | Design Decisions | 6 |
| 4 | Framework and Tools | 6 |
| 4.1 | Gathering the Data | 6 |
| 4.1.1 | Common Crawl | 6 |
| 4.1.2 | Other Data Sources | 7 |
| 4.2 | Selecting Relevant Documents | 7 |
| 4.3 | Extracting Relations from Documents | 8 |
| 4.3.1 | Defining Classes | 8 |
| 4.3.2 | Pre-processing | 8 |
| 4.3.3 | training | 9 |
| 4.3.4 | classifying | 9 |
| 4.4 | Storing and Ingesting the Data | 10 |
| 4.4.1 | Graph Database, Search Engine or Traditional Database | 10 |
| 4.4.2 | Comparing Graph Databases | 10 |
| 4.4.3 | Using Neo4j for Storage and Ingestion | 11 |
| 4.5 | Interacting with the Data | 12 |
| 4.5.1 | Query Language | 12 |
| 4.5.2 | Query Composer Interface | 12 |
| 4.5.3 | Interactive Search | 12 |
| 4.5.4 | Conclusion | 12 |
| 4.6 | Visualising the Data | 13 |
| 4.6.1 | Representing the Data Graphically | 13 |
| 4.6.2 | Using Geographical Maps | 13 |
| 4.6.3 | Handling Map Clutter | 13 |
| 5 | Validation | 14 |
| 5.1 | Tests | 14 |
| 5.1.1 | unit tests | 14 |
| 5.1.2 | Integration tests | 14 |
| 5.1.3 | System tests | 15 |
| 5.1.4 | Acceptance Testing | 15 |
| 5.2 | SIG | 15 |
| 6 | Conclusion | 15 |
| A | Better Code Hub Guidelines | 16 |

1 Introduction

With the development of future cities in mind, the interest in city networks has grown over the years. However, appropriate information on how cities are connected and the strength of these connections is hard to find. According to Short et al. comparative statistics are not easily available and common assertions are repeated [23]. Although more research was published since then, for example the work of DeRudder et al. [6], there is still a lot of information missing on relationships between cities. In order to get more insight into these relationships, there is a need for complex data analysis.

The huge amount of textual data generated online and the numerous historic archives, such as for example Delpher¹ and the British Newspaper Archive², are great sources of information on social and economic behaviours. The clients hypothesis is that advanced text mining on newspapers and web pages containing city names would allow for a better understanding of the role of information in shaping urban systems. Similar to research efforts in other domains, such as financial trade [21] and sales forecasting [29], the clients wish is to develop an application that captures urban-urban interactions. These interactions are retrieved from information corpora through the co-occurrence of geographical names in textual data. An example of how one could try to achieve this using the Google search engine³ is "Rotterdam + Amsterdam " OR "Amsterdam + Rotterdam", which searches for the co-occurrence of Amsterdam and Rotterdam. However, manually processing all results a search engine yields is not feasible because someone would have to read each page to determine which types of relationships the page contains. An application should process all the pages that contain co-occurrences of cities to determine what type of relations, for example transportation or leisure, between cities can be extracted from the document. Thus, we will answer the following question: how can the strength of relationships between cities be extracted and visualised from open data using software?

First, we discuss related work in section 2. Second, we identify the requirements for a solution to the problem and discuss issues that might arise in section 3. Third, we develop a methodology for a framework that satisfies the requirements and tackles the issues in section 4. In the fifth section we discuss evaluating the system. We conclude in section 6 with the results of our research.

2 Related Work

In the 1960s, progress in the knowledge of urban networks was made feasible by the rapidly increasing availability of useful data [28].

One of the most common methods used is the interlocking network model (INM) [27]. This model assumes cities have a flow of knowledge connection if there are offices of the same company in those cities. The biggest problem with this is that it is very limited. It only includes one relation type and it is disputable whether this is a good measurement for the relation.

Since one decade, we are facing a new revolution in data production and processing. Developing innovative methods to extract relevant information from Big Data could again allow important progress in the knowledge of the interurban relationships.

When looking at digital data there are two different approaches: the cyberspace and the cyberplace [7]. The cyberplace measures relations by using the infrastructure of the internet, for example internet hyperlinks, the structure of search engines and email traffic. **TODO: The bulk of efforts have concentrated on this “backbone” of Internet made of cables and routers ([3,10]**

The cyberspace method focuses on the virtual communication of people through connected devices. One approach is by registering and mapping the number of pages indexed by search engines for queries containing the names of two cities **TODO: Devriendt et al., 2008; Janc, 2015).. TODO: In 2010 Brunn et al. registered the number of pages indexed for two cities followed by key words such as “global financial crisis” or “climate change” to evaluate the linkages between global cities concerning these two subjects. More advanced textual analysis on websites and search engine queries containing co-occurrences of city**

¹<http://www.delpher.nl>

²<http://www.britishnewspaperarchive.co.uk/>

³<https://www.google.com>

names would allow to better characterize digital links between cities in a more systematic approach.

Therefore we intend develop a computer program which does this. It will automatically find predefined relations between cities and their strength by using all pages available from search engines. In the following chapter we will investigate the requirements such a program should have.

3 Requirement Analysis

In this section, we first present user stories that were created together with the client. Next, we define the design goals. Then, we list the requirements which followed from the user stories and which the application should meet. To do so, we use the MoSCoW method [4] as a prioritisation technique. Lastly, we discuss the design decisions that follow from the design goals and the requirements.

3.1 User Stories

Together with the client, several user stories are identified and listed below:

1. As a user, I want to be able to see all the identified relations between all cities, so that I can reason about interesting patterns.
2. As a user, I want to be able to access the processed data in an excel file. I want this to be available per relation type and as a total of all relations, so that I can apply my own models on the data when I want to.
3. As a user, I want to be able to see relation strengths. Which can be expressed by counting the relations
4. As a user, I want to be able to (de)select cities in the user interface, so that I can create a network of cities connected with relations. Where a network of cities will be the cities on a map and the different types of relationships such as Transportation will be the connecting edges between cities.
5. As a user, I want to be able to (de)select relations between cities in the user interface, so that I can inspect only the relations I am interested in at that time. For example, as a user I might only be interested in the Transportation relationship between Amsterdam and Rotterdam.
6. As a user, I want to be able to change the colours associated with the different relation types, so that I can adjust the styling to my own preferences.
7. As a user, I want to be able to export an image of the map that I composed in the user interface so that I can use it for presentations, papers or educational purposes

3.2 Design Goals

The high-level design goals for this project have been provided by the client. These serve as a guideline to determine the priority label of the specific requirements as defined in section 3.3. The design goals are listed below, ordered by priority.

credible The results of the project will be used in research on intercity relations. Therefore, the results must be reliable and verifiable. This means that the application should produce the same results given the same input and it should be possible to manually access the input to verify the output of the application.

understandable The results of the application should be visually understandable, in order to make it easy for the client to deduce conclusions.

scalable During the project a TU Delft server will be used with a limited amount of resources. Therefore only .nl pages will be used as input to limit the amount of data storage and processing power needed. However, allowing for investigating other domains would greatly help the client in a later stadium, which means that the system would have to be scalable where possible. For example, using a dedicated database which can be spread across clusters.

pluggable It might be interesting for the user to let the application perform analysis on different data sets without the need of a developer. So if possible within the time constraints the application should be able to use any form of textual input data.

exportable Besides making the results available visually, all the relevant numeric data should also be exportable, for example as a matrix.

fast development Because of the time constraints of the project we need a fast development cycle. As a result of that, choices regarding tools, applications and programming languages are to be made with the time constraint taken into account.

3.3 Product Requirements

As mentioned in the introduction of section 3 we will be using the MoSCoW method prioritisation technique. Four levels of priority are defined: must have, should have, could have and won't have (also known as would like). We also differentiate between functional and non-functional requirements.

3.3.1 Must Have

Requirements labelled as "must have" are key to the minimal performance of the application. If they are not met, the application can be considered a failure.

1. Data that is of relevance for the UrbanSearch project, should be mined from the Common Crawl web corpus (see section 4.1.1) and stored for further processing/access.
2. There has to be a way to extract the collected data based on relations between cities/nodes.
3. A machine learning algorithm should analyse the collected data and label different types of relations that are important for intercity relations.
4. A front-end should be built for the project. This front-end should visualise basic relations and statistics and can be used for presentations and educational purposes.
5. Several statistically important aspects of intercity relations should be extracted from the data set. These statistics should be easily accessible and visualised to the end user. Furthermore, it should be easy to extend or update the list of statistics that are associated with a relation.

3.3.2 Should Have

"Should have" requirements are those that greatly improve system performance and/or usability but might not fit in the available development time.

1. Relations between cities should be accessible hierarchically. This means that there is the possibility to explore a relation and, provided that this relation has sub-types associated with it, the relation can be expanded in the different sub-types of the relation.
2. It should be possible to train the machine-learning algorithm by feeding it labeled data.
3. It should be possible to add data sets with more than 1 million documents on which the system can perform its data mining routines. This way a data set can be created that contains potentially interesting information for intercity relations.
4. The application should be able to deal with the fact that the same city can have different names in different languages/dialects. It should still be able to extract and group relevant data correctly (e.g. 'The Hague' and 'Den Haag' should be viewed as the same city).

3.3.3 Could Have

Requirements labelled as "could have" are useful and should be included in the system if time and resources permit.

1. The system should use Delpher to characterise relationships between a region and cities outside that region based on newspapers in those regions. For example, the local newspaper of the province Gelderland writing about the city of Alkmaar. These relationships are either simple or complex information flows **TODO: Definition of Information flow**. According to our client simple information flows consist of a newspaper located in city i publishing about city j , whereas complex information flows are co-occurrences of cities in an article.
2. The relations that are extracted from the data by the machine learning algorithm have to be visualised in a way that makes it easy to compare the different relations for the end user. For example, a split-screen comparison in the user interface or an export of graphs comparing selected relations.

3.3.4 Would Like

"Would like" requirements have been agreed upon to be not important to include within the current time schedule. However, they can be included in future releases.

1. The application would be able to show all connections of all places on the map at the same time.
2. Using data from top-level domains other than `.nl`.

3.4 Design Decisions

To be able to have a fast development cycle and leverage our experience we chose to develop the application using Python. We plan to not only test the code we deliver thoroughly, but also to cross-validate the obtained results **TODO: Refer to section with evaluation protocol**.

4 Framework and Tools

In this section, we gradually develop the framework. First, we decide what data source to use initially and describe how to use the data. Second, we present the method of choice to perform initial data filtering. Third, we agree on which data storage and ingestion to use. Fourth, we provide a methodology to group the data. Fifth, we decide how users will interact with the obtained results. Last, we select the visualisation tools to be used. **TODO: Afbeelding van hoe systeem werkt qua inputs/outputs en stappen**

4.1 Gathering the Data

As explained in section 3.2, data sources should be pluggable. An initial corpus of documents is needed to base the project, which we will decide on in this section. Nowadays many people have access to the internet and for a lot of people the internet is probably also their primary source of information. Next to that, the internet also contains vast amounts of documents which could shed some light on relations between cities. Therefore, the decision was made to use web-data as a data source. To avoid duplicate work, which would mean crawling the web, a logical choice is to use Common Crawl as a data source.

4.1.1 Common Crawl

Common Crawl [5] is a freely accessible corpus of pages across the web, updated and released on a monthly basis. Many researchers have used the data for various purposes [16, 24, 25]. Since the project requires analysis on a very large set of documents, the corpus is a very suitable candidate for us to work with.

The data from Common Crawl comes in three formats⁴:

⁴<https://gist.github.com/Smerity/e750f0ef0ab9aa366558>

WARC This is the default and most verbose format. It stores the HTTP-response, information about the request and meta-data on the crawl process itself. The content is stored as HTML-content.

WAT Files of this type contain meta-data, such as link addresses, about the WARC-records. This meta-data is computed for each of the three types of records (meta-data, request, and response). The textual content of the page is not present in this format.

WET This format only contains extracted plain text. No HTML-tags are present in this text. For our purposes, this is the most useful format.

Common Crawl stores these pages in the following way: each archive is split into many segments, with each segment representing a directory. Every directory contains a document listing file and a folder for each file format (WARC, WAT and WET), which in turn contains the compressed pages belonging to the segment. To be able to efficiently get a single page, Common Crawl indexes the segments to directly map URLs to document locations using an offset and length which can be found using the Common Crawl index⁵. Since WAT- and WET-files can be generated from WARC-files, they only provide such indices for WARC-files. If no file index is provided with a data request, an aggregated compressed file of all files of the requested format is returned.

For extracting data from Common Crawl, many open-source libraries are available. Common Crawl's official website refers to `cdx-index-client`⁶ as a command line interface to their data indices. It allows for, among others, specifying which data set to use, supports multiple output formats (plain text, gzip or JSON) and can run in parallel. Since this library only retrieves the file indices, we need another way to actually retrieve the pages pointed to. However, there is a problem with this: we are only interested in WET-files, but Common Crawl does not have WET-files indexed. We would therefore have to collect the WARC-files and convert them to WET-files ourselves, requiring us to parse HTML for every document we are interested in.

As mentioned in the design goals section not all available web-data will be used due to limited resources. A simple query `url=*.nl&output=json&showNumPages=true` on the CC-MAIN-2017-13 index using the online interface⁷ yields 1676 pages. Pages in this sense are listings of 15000 indices, so there are roughly 25 million entries in total out of the 2.94 billion pages available in Common Crawl. It is very important to note that searching for a top level domain like `.nl` only includes the first page of every matching domain. To get all pages, additional queries for each site with more than one page are to be performed.

4.1.2 Other Data Sources

Besides Common Crawl, there are a plethora of other sources that might contain valuable information. The most notable is the Dutch royal library, Delpher⁸. It contains millions of Dutch digitalised newspapers, books and magazines from the fifteenth century up until about 1995. Because of this, it is a useful resource for historical research. Additionally, Statistics Netherlands⁹ is the governmental organisation collecting statistical data about the Netherlands and comes with an API, making most of their data publicly accessible. The NOW Corpus¹⁰ collects newspaper and magazine articles through Google News and provides several tools to perform queries on this data. It can also be downloaded.

Due to time and resource constraints, we have chosen to exclude these from the project. Of course, in future versions, other data sources could be included.

4.2 Selecting Relevant Documents

Because not all data from information sources such as Common Crawl is relevant to find relationships between cities, the data needs to be filtered. One way to do this, is to only select the data that mentions at least two different cities. Because the data is plain text, we need a way to scan through the text and determine if the text indeed has a co-occurrence of two different cities. Making use of the comparative analysis of Rasool et al. [22], we chose the Aho-Corasick algorithm [1], which is a multi-pattern exact string matching algorithm. Using this algorithm, a predefined list of cities can be matched against the

⁵<http://index.commoncrawl.org>

⁶<https://github.com/ikreymer/cdx-index-client>

⁷http://index.commoncrawl.org/CC-MAIN-2017-13-index?url=*.nl&output=json&showNumPages=true

⁸<http://delpher.nl>

⁹<https://www.cbs.nl/en-gb>

¹⁰<http://corpus.byu.edu/now/>

text of a web page or document. If at least two cities from the list appear in the text, we mark it as a useful document.

The decision to use the Aho-Corasick algorithm is strengthened by the fact that a well documented and stable Python library exists, which implements the aforementioned algorithm. This library is called `pyahocorasick`¹¹ and is a fast and memory efficient implementation of the Aho-Corasick algorithm.

We make a selection of documents without storing the documents first, because storing and indexing all documents is not feasible due to storage constraints. For the .nl web pages only would need about 250 GB of storage and to store all available documents around 250 TB of storage would be needed. As we do not have access to a fast and large data storage platform, we will not store and index everything first and then delete irrelevant documents. However, to test if finding and storing relationships between cities is fast enough when the documents are actually stored on disk a random selection of 1 million documents will be downloaded. Processing the already stored documents could finish within one day whereas downloading all documents will most certainly take multiple days.

4.3 Extracting Relations from Documents

Now that a selection of relevant documents has been made, we can make an attempt to identify the relations between cities based on these documents. Since labelling every relevant document by hand is not feasible an automated approach is desirable. One way to automate this process is by identifying intercity relations using machine learning. Machine learning algorithms can be roughly divided into two distinct groups: Supervised and unsupervised algorithms. Supervised algorithms expect an input set and a corresponding output set by which an model is trained to predict unseen instances of the problem. Unsupervised algorithms identify clusters of entities based on similarities in the feature set corresponding to said entity. Considering the fact that we have a strict time schedule of only two months to develop the complete system, we decided with our client to go with the supervised approach. The main reason for this choice is that the training and tweaking of supervised algorithms can be done faster compared to unsupervised algorithms. The main reason for this is that we do not need the complete data-set to start training a supervised model, while for the unsupervised case the complete set is needed.

4.3.1 Defining Classes

Our choice of using classification has naturally lead to the need for categories we want to identify within the collected documents. Together with our clients, Dr. Evert Meijers and Antoine Peris, we identified the following categories which are useful to identify from the collected documents:

1. Commuting
2. Leisure
3. Residential mobility
4. Education
5. Collaboration
6. Transportation
7. Other

These categories represent topics that are of interest for our clients. They relate to research that is being done by our clients and to relations that were deemed important in previous research on intercity relations.

4.3.2 Pre-processing

For text pre-processing there are a number of tools available. We used NLTK [?] removing stopwords and Scikit-Learn [14] for Modelling.

¹¹<https://pypi.python.org/pypi/pyahocorasick/>

1. Tokenization.
splitting up the text into words and other symbols called tokens. For this we implemented the tokenizer NLTK provides. Testing showed that it does what it is supposed to do.
2. Removing stop words.
Removing all common words (the, a, an etc) and symbols ('.', ',', '!', etc). For removing stopwords we used a list from NLTK containing dutch stopwords. For removing symbols we made a method that removes all non alphabet values except spaces from text.
3. Stemming.
Reducing derived words to their stem (e.g. fishing -> fish). We first wanted to use NLTK, which itself implements several stemmers including SnowBallStemmer [18] and Porter [19]. However research from the Rijks Universiteit Groningen (RUG) by Tanja Gaustad and Gosse Bouma [9] evaluated a Bayesian text classification system with either no stemming or the Porter or dictionary based stemmer. Which concluded stemming does not lead to significant change in classification accuracy.
4. Modelling (TF-IDF).
Before we can enter this tokenized data into the machine learning algorithm we first need to transform it to the correct model. Gensim provides two ready to use method for this to change tokenized text into the proper input data. The first is token2ID which gives each unique word a numeric ID and counts the amount that those words occur in a text and the second is tfidf which changes the amount that words occur in tf-idf values. tf-idf means term frequency inverse document frequency. Term frequency is the number of times the word occurs in a document and idf is the number of documents divided by the number of documents where the term occurs. These two values are then multiplied. The results are normalised to make up for different document lengths [https://radimrehurek.com/gensim/models/tfidfmodel.html].

4.3.3 training

In order to use supervised machine learning, we first need to find training data for each class. We thought about two ways to do this. In order to get a good training sample for these categories, we discussed two different methods.

The first option was to query for results from news(paper) sites. However, this approach did not give us the desired results. The reason for this is that the categories we find in the newspapers do not match the categories Antoine and Evert provided us with. We also tried using the search engine from those news sites, however using those we found completely unrelated articles to the search queries (A search query about 'verhuizen' resulted in an article about parrots and minecraft).

The second option we are considering is to use Google Custom Search to obtain results from Google using the categories/keywords Antoine and Evert provided us with. A quick test with this last method provided us with quite good results (although we still see some noise in the results this is less than with the search engines from news sites).

4.3.4 classifying

For classifying there are a multitude of algorithms available. For choosing the classifier we make use of the Microsoft Azure Machine Learning Test Sheet [15]. Several factors should be taken into account when choosing an algorithm. These are:

1. Accuracy - How well the algorithm separates the websites.
2. Training Time - How long it takes to train the algorithm.
3. Linearity - Linear regression assumes data trends follow a straight line. This is trade-off between accuracy and speed.
4. Number of Parameters - Adjustable parameters increase the flexibility of the algorithms. This is a trade-off between training time and accuracy.

5. Number of Features - A large number of features can make some algorithms unfeasibly long. Especially text data (what we are using!) has a large number features. Support Vector Machines are especially well suited in this case.
6. Special Cases - Some learning algorithms make particular assumptions about the data or the results.

For textual data especially support vector machines are recommended, so it is most likely we will choose that machine learning algorithm. Depending on whether we have time we might do some tests before making our decision however. The sci-kit package [14] provides an easy to implement module to use this.

4.4 Storing and Ingesting the Data

In this section we will discuss which data storage solution we are going to use and why. We will compare a few options and select the best. We will then briefly explain how it works and how we plan to use it.

4.4.1 Graph Database, Search Engine or Traditional Database

Now that relationships have been extracted from the documents and web pages, we need to store these. Note that is not necessary to store all the documents, because the relationships which have been extracted are the information that is actually worth storing. Therefore, the original document can be discarded. However, the index pointing to the original document will be stored to be able to view it on demand. To store these relationships and indices there are two possibilities: (1) graph databases and (2) traditional relational databases.

Because visualisation of the network of cities as a graph is an important part of the application, and relations between cities play a key role in the system, we need a database that is designed for these features. Relations are the most important in the graph data model, where this is not true for traditional relational databases. **TODO: Meer argumentatie waarom Graph DB beter** Therefore, we are confident that a graph database is the best choice.

4.4.2 Comparing Graph Databases

Next, the type of database needs to be selected. For this, six of the most popular databases according to the solid IT Graph DBMS ranking¹² are rated on five important aspects. These are, is the graph database open-source, scalable and free and does it support Python and have built-in visualisation.

| name | Open-source | Scalable | Python support | Free | Built-in Visualisation |
|----------------|-------------|----------|----------------|----------------|------------------------|
| AllegroGraph | ✗ | ✓ | ✓ | ✗ ^a | ✗ ^b |
| ArangoDB | ✓ | ✓ | ✓ | ✓ | ✓ |
| Neo4j | ✓ | ✓ | ✓ | ✓ ^c | ✓ |
| OrientDB | ✓ | ✓ | ✓ | ✓ | ✓ |
| Teradata Aster | ✗ | ✓ | ✓ | ✗ | ✗ ^d |
| Titan | ✓ | ✓ | ✗ | ✓ | ✗ ^e |

^a Only free up to 5 million triples

^b With separate tool called Gruff: <https://allegrograph.com/gruff2/>

^c Non-commercial use

^d Using a separate tool Aster AppCenter

^e Using a separate tool

From this table, it can be deduced that three of these graph databases are viable candidates: ArangoDB, Neo4j and OrientDB. For this project, Neo4j is the best choice because of three reasons. Firstly because we have experience with Neo4j, which means less time will be spent on getting to know the graph database and functionality. Secondly because it is by far the most popular graph database¹³. Thirdly, since Neo4j is the most popular graph database, the support community and amount of available examples is large.

¹²<https://db-engines.com/en/ranking/graph+dbms>

¹³<https://db-engines.com/en/ranking/graph+dbms>

Labeled Property Graph Data Model




Figure 1: The Neo4j model

4.4.3 Using Neo4j for Storage and Ingestion

Neo4j is a highly scalable native graph database that leverages data relationships as first-class entities [17]. It is the single highly scalable, fast and ACID compliant graph database available. ACID stands for the four properties atomicity, consistency, isolation and durability of transactions in database systems that ensure reliability for query results [11]. Additionally, Neo4j is free to use for non-commercial purposes. To illustrate how scalable Neo4j is, consider that very large companies such as eBay, Cisco, Walmart, HP and LinkedIn¹⁴ use it in their mission-critical systems. Holzschuher and Peinl compared the performance of Neo4j to the more classic and commonly used NoSQL databases and found that the more natural representation of relationships resulted in significant performance increase gains [12].

There are some specific aspects of Neo4j that make it a very suitable candidate for this application. These are:

properties Any entity in the Neo4j graph can be given properties (key-value pairs) containing information about the entity. Properties are primarily meant to provide additional information and are less suitable to be queried on. As an example, a city can have a number of inhabitants and districts attached to it as a property.

labels Nodes can be tagged with a label, describing their roles in the network. These annotations are especially useful to filter the data set on one or more categories. For example, a city can be labelled as "capital" to be able to distinguish between regular and capital cities.

relations Nodes can be connected using relationships. These are always directed, typed and named and can have properties. Using these properties, one can control how the graph is traversed. For example, if a path (relationship) is to be avoided unless absolutely necessary, the relation can be given a high cost. To give importance to some relationship, one could also assign a strength score to it. Since relationships are handled efficiently by Neo4j, nodes can have any number of relationships linked to it without compromising performance. For our purpose, a relation could comprise the strength of the relationship between two cities (nodes).

The Neo4j model can be depicted as shown in figure 1. It consists of nodes, relationships (edges), properties (within the nodes) and labels (rectangular blocks above the nodes).

Besides the aforementioned useful properties of Neo4j, the graph can be put to good use for visualising the global urban network. By adding a location property to a city, nodes and relations can be mapped directly to a geographical map. Most importantly, indices of text files can be stored that mention the city as properties of nodes. That way, we are able to generate a subset of files that can be analysed for calculating the strength of the relationship between the nodes.

¹⁴<https://neo4j.com/customers/>

4.5 Interacting with the Data

After having filtered and classified the data, the framework should provide a means for the client to interact with the resulting data. In this section, several ways to do so are compared, after which we decide which path to take.

4.5.1 Query Language

One possibility is to let the client query the data. For this, we designed a simple, easy to use query language specific to the domain of research. It has the following syntax:

| | |
|-----------|-----------------------------------|
| ! | Logical NOT operation |
| & | Logical AND operation |
| | Logical OR operation |
| (A & B) | Grouping of clauses |
| A > R > B | Relation R between cities A and B |

In figure 2, an example is shown that queries the "Shopping" relation between Rotterdam and Amsterdam and between Rotterdam and Den Haag.

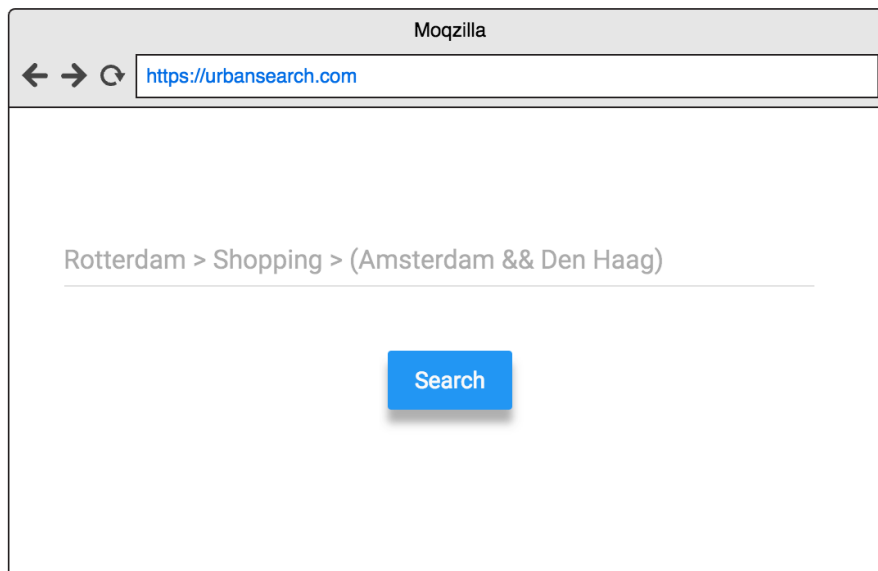


Figure 2: Example interface for the query language

4.5.2 Query Composer Interface

Another possibility is to offer the client a query composition interface. This interface would have the same functionality as the previously mentioned query language, but is more intuitive to use for new users. An example of the interface is given in figure 3.

4.5.3 Interactive Search

The last option we investigated is an interactive approach to querying data. For this, the client interacts with a map containing relations and cities. A very simple example is given in figure 4.

In this setup the user clicks on cities and relations on the map. This event triggers a query on the back-end and the resulting data is visualised on the map. An example of such an event is to show information about the selected city.

4.5.4 Conclusion

In association with the client, we conclude that the best option to go with is the interactive map. This way, the client has easy access to the data and this pattern of interaction best suits the work flow that




Figure 3: Example interface for the query composer

the client envisioned prior to the project. The creation of complex queries is also made a lot easier. The user does not have to write or compose a complex query in advance but can do it directly on the map. Thus, retrieving a visual representation of several cities, interconnected with multiple relations, only involves selecting cities and relations on the map. Interaction directly with the map also reduces the need to go to a separate page to compose a query. This speeds up the use of the system by reducing page loads and it interrupts the work flow of the user less.

4.6 Visualising the Data

This section focuses on the visual representation of the processed data. This should be done in way that is easy to comprehend for users and helps them to interpret the data. To reach these goals, we identified the clients needs and desires. We have discussed the preferences of the client and have drawn up a global plan, which we present below.

4.6.1 Representing the Data Graphically

Since we are dealing with strongly related data, it is a natural choice to represent the data as a graph. The choice was made, in association with the client, to show the nodes and relations on a geographical map. Visualising cities on a map is intuitive to the user and we think this will increase the ability of users to interpret the information in a productive manner.

4.6.2 Using Geographical Maps

We investigated two map libraries we can use to display our data on a map. The first one is Google Maps, which can be used freely and offers a lot of customisation options. The API is well defined and some of the group members have previously worked with it. The second option we investigated is Leaflet. Leaflet is an open-source JavaScript library that provides responsive maps. It also has fine grained API and lots of plugins available. Both libraries are well suited for our needs. However, we decided to go with Google Maps, because of the existing experience of the group members. Also, we feel Google Maps is better supported through documentation and communities.

4.6.3 Handling Map Clutter

One of the challenges of visualising networks, as stated in [2], is the occurrence of so-called map clutter. Map clutter means the network is displayed as an incomprehensible set of nodes and edges. Several methods to prevent this are given in [2]. We will adopt some of these methods in our application, as explained next.



Figure 4: Example of an interactive map

Users should be able to select what information they want to display. This will be included in the system by allowing the user to select cities and relations, enabling them to filter nodes and edges. The use of different sizes for nodes and edges or other attributes that are displayed can convey extra information to the user. We will use this to represent, for example, city population and exact strengths of relations. The use of colour is another method mentioned in [2]. We will use colours to represent different types of relations and utilise colour intensity and opacity to represent the strengths of these different types of relations.

5 Validation

5.1 Tests

We will test the program using four different testing methods. The first is unit testing, which tests the separate components individually. Integration testing for testing how well different components work together. system testing for testing the different entire system and Acceptance testing for testing how well the clients think the program works.

5.1.1 unit tests

Unit testing is done by writing automatic tests and making sure they pass every time the tests are executed. Unit tests test each method of a function separately, checking that the method does what it is supposed to do. If the method would need information from outside the class that information is mocked. This means that instead of using that other class, a fake object is made which returns a fake value. This ensures the tests will never fail due to changes in other classes.

selenium tests

Selenium tests are automated tests that are run to test the front-end of systems. Just as normal unit testing, selenium testing is useful for regression testing.

5.1.2 Integration tests

Integration testing uses automated tests which test how well different components of the system work together. This is done more or less the same as unit testing, however whilst you would mock methods from other classes in unit testing, with integration testing you do not. It is assumed that the separate modules are unit tested, therefore if an error occurs it is because something is wrong with the interaction between the modules and not with the modules themselves.

5.1.3 System tests

We are also planning to use system testing. System testing provides a more complete test of the entire system. This means it is useful to detect faults in the overall system, but less easy to determine where these faults may be located. System testing is done manually, which means the tests can not be easily repeated when the system changes whilst with other testing techniques this is possible.

5.1.4 Acceptance Testing

Last we will also use acceptance testing. This is testing done to see if the software does what the clients are expecting it to do. These tests are therefor also executed by the clients manually. Afterwards they can say what worked, what didn't worked, what was missing and what could be improved. Part of the acceptance test will be to see how well the machine learning classifier is. Since we need documents fitting to the different classes this is part of acceptance testing too.

machine learning

To test the machine learning we will manually search for a small set of pages per class. We will then manually decide to which class we would expect those pages to be sorted. Then we will run the algorithm and check whether or not a page was correctly sorted. There are 4 outcomes per page per class: true/false positives and true/false negatives. Positives are the pages that should be sorted to that class; true means they were sorted to that class, false meaning they were not sorted to that class. Negatives are pages that should not be sorted to that class and in this case true means they were indeed not sorted to that class and false means they were. So the true values are the ones that should be high. From this we can say how accurate our classifier is.

5.2 SIG

SIG [26], short for software improvement development group, is a tool to give deep code visibility we require so we can remedy current problems and prevent future vulnerabilities. SIG includes Better Code Hub [13] which checks our code according to 10 guidelines as can be seen in appendix A. The great thing about Better Code Hub is that it can be run at anytime. We can check better code hub whenever, whilst for SIG we have to send in our code and wait for feedback.

6 Conclusion

First, we discussed related work. We saw that there are currently two methods for analysing the relations between cities. Manually analysing search engine data is very slow and requires a lot of man-hours and looking at the different locations where businesses are located is only interesting for the economic relation and still misses a lot of data.

Second, we identified the requirements for a solution to the problem and discuss issues that might arise. The used the MoSCoW model to describe the importance of the different requirements. The most import must haves we found are being able to input place names, displaying a map with the connection data and being able to extract this data.

Third, we developed a methodology for a framework that satisfies the requirements and tackles the issues. We decided to start by using data from Common Crawl, although we might later extend this to other data sources such as Delpher. After selecting relevant data (data which contains 2 or more city names) we store the data with Neo4j. We then use clustering and classifying machine learning to group the data. First we use this on all data to get the general groups (e.g. economy, health-care, immigration) and then we use this on the data per pair of cities to see what the important connection types for each city are. Then we link these connections to the general groups ('fish' might relate to economy.. etc). To visualise this data we use the graph Neo4j provides.

Last we discussed how to validate the methodology applied. We will test our code with four levels of testing and send our code to SIG for analysis. We will use unit tests to tests individual modules, integration tests to check how different modules that call upon each other work together, system tests to check the entire system ourselves and acceptance testing to see how well the system fits the clients

needs. To ensure visibility and maintainability in our code send the code twice to SIG for feedback and frequently use the service of Better Code Hub.

With this setup we should be able to make a well-tested functioning system that meets the requirements of our clients. Furthermore, using this system will enable us to answer the question "how can the strength of relationships between cities be extracted and visualised from open data?"

A Better Code Hub Guidelines

Better Code Hub [13] checks our code according to ten guidelines:

1. **Write short units of code**
Units of code should be no longer than 15 lines.
2. **Write simple units of code**
Separate units of code should contain no more than 4 branch points (if, for, while, etc)
3. **Write code once**
Shared code should be extracted, either to a new unit or to a super class
4. **Keep unit interfaces small**
The number of parameters per unit of code should be no more than four.
5. **Separate concerns in modules**
Identify and extract responsibilities of large modules to separate modules and hide implementation details behind interfaces.
6. **Couple architecture components loosely**
minimizing the amount of interface code (e.g. by using 'abstract factory' design pattern)
7. **Keep architecture components balanced**
Organize code in such a way that the number of components is between 2 and 12, and ensure the components are of approximately equal size (keep component size uniformity less than 0.71).
8. **Keep your codebase small**
Refactor existing code to achieve the same functionality using less volume, and prefer libraries and frameworks over "homegrown" implementations of standard functionality.
9. **Automate tests**
Add tests for existing code every time you change it.
10. **Write clean code**
Remove useless comments, commented code blocks, and dead code. Refactor poorly handled exceptions, magic constants, and poorly named units or variables.

References

- [1] Alfred V. Aho and Margaret J. Corasick. Efficient string matching: An aid to bibliographic search. Commun. ACM, 18(6):333–340, June 1975.
- [2] R. A. Becker, S. G. Eick, and A. R. Wilks. Visualizing network data. IEEE Transactions on Visualization and Computer Graphics, 1(1):16–28, Mar 1995.
- [3] Junho H Choi, George A Barnett, and BUM-SOO CHON. Comparing world city networks: a network analysis of internet backbone and air transport intercity linkages. Global Networks, 6(1):81–99, 2006.
- [4] Dai Clegg and Richard Barker. Case method fast-track: a RAD approach. Addison-Wesley Longman Publishing Co., Inc., 1994.
- [5] Common Crawl. Common crawl. <https://commoncrawl.org/>, 2017. Accessed: 2017-04-25.

- [6] Ben Derudder and Frank Witlox. An appraisal of the use of airline data in assessing the world city network: a research note on data. Urban Studies, 42(13):2371–2388, 2005.
- [7] Lomme Devriendt, Ben Derudder, and Frank Witlox. Cyberplace and cyberspace: two approaches to analyzing digital intercity linkages. Journal of Urban Technology, 15(2):5–32, 2008.
- [8] Elasticsearch. Elasticsearch. <https://www.elastic.co/products/elasticsearch>. Accessed: 2017-04-26.
- [9] Tanja Gaustad and Gosse Bouma. Accurate stemming of dutch for text classification. Language and Computers, 45(1):104–117, 2002.
- [10] Sean P Gorman and Edward J Malecki. The networks of the internet: an analysis of provider networks in the usa. Telecommunications Policy, 24(2):113–134, 2000.
- [11] Theo Haerder and Andreas Reuter. Principles of transaction-oriented database recovery. ACM Computing Surveys (CSUR), 15(4):287–317, 1983.
- [12] Florian Holzschuher and René Peinl. Performance of graph query languages: comparison of cypher, gremlin and native access in neo4j. In Proceedings of the Joint EDBT/ICDT 2013 Workshops, pages 195–204. ACM, 2013.
- [13] Software improvement group. better code hub. <https://bettercodehub.com/>, 2017.
- [14] Scikit learn developers. Scikit-learn. <http://scikit-learn.org/stable/index.html>, 2017. Accessed: 2017-04-26].
- [15] Microsoft. Microsoft azure machine learning algorithm cheat sheet. <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>, 2017. Accessed: 2017-04-28.
- [16] Hannes Mühleisen and Christian Bizer. Web data commons-extracting structured data from two large web corpora. LDOW, 937:133–145, 2012.
- [17] Neo4j. Neo4j, the world’s leading graph database. <https://www.neo4j.com>. Accessed: 2017-04-26.
- [18] Richard Boulton Olly Betts. Snowball. <http://snowball.tartarus.org/algorithms/dutch/stemmer.html>, 2017.
- [19] M.F. Porter. Porter. <https://tartarus.org/martin/PorterStemmer/>, 1980.
- [20] Michael E Porter. Location, competition, and economic development: Local clusters in a global economy. Economic development quarterly, 14(1):15–34, 2000.
- [21] Tobias Preis, Helen Susannah Moat, and H Eugene Stanley. Quantifying trading behavior in financial markets using google trends. Nature: scientific reports, 2013.
- [22] Akhtar Rasool, Amrita Tiwari, Gunjan Singla, and Nilay Khare. String matching methodologies: A comparative analysis. REM (Text), 234567(11):3, 2012.
- [23] John Rennie Short, Y Kim, Merje Kuus, and Heather Wells. The dirty little secret of world cities research: data problems in comparative analysis. International Journal of Urban and Regional Research, 20(4):697–717, 1996.
- [24] Sameer Singh, Amarnag Subramanya, Fernando Pereira, and Andrew McCallum. Wikilinks: A large-scale cross-document coreference corpus labeled via links to wikipedia. University of Massachusetts, Amherst, Tech. Rep. UM-CS-2012-015, 2012.
- [25] Jason R Smith, Herve Saint-Amand, Magdalena Plamada, Philipp Koehn, Chris Callison-Burch, and Adam Lopez. Dirt cheap web-scale parallel text from the common crawl. In ACL (1), pages 1374–1383, 2013.
- [26] software improvement group. Sig. <https://www.sig.eu/>, 2017.
- [27] Peter J Taylor. The interlocking network model. International handbook of globalization and world cities, pages 51–63, 2012.

- [28] Gunnar Törnqvist. Flows of information and the location of economic activities. Geografiska Annaler. Series B, Human Geography, 50(1):99–107, 1968.
- [29] Lynn Wu and Erik Brynjolfsson. The future of prediction: How google searches foreshadow housing prices and sales. In Economic analysis of the digital economy, pages 89–118. University of Chicago Press, 2014.